

Screenshot of an AWS S3 console interface showing the contents of the 'incoming' folder in the 'student-data-lake-34' bucket.

The browser tabs at the top include: Course Inv, Launch AW, Query edit, s3-daily-re, CloudWatch, student-d, Subscriptions, Redshift, and several others.

The URL in the address bar is: us-east-1.console.aws.amazon.com/s3/buckets/student-data-lake-34?region=us-east-1&prefix=incoming/&showversions=false

The AWS navigation bar shows: Account ID: 5673-2992-3344, United States (N. Virginia), and the user's email: vorlabs/user4548998=rohamrutkar2525@gmail.com

The breadcrumb navigation path is: Amazon S3 > Buckets > student-data-lake-34 > incoming/

The main content area displays the 'incoming/' folder with one object listed:

Name	Type	Last modified	Size	Storage class
pipeline_data_quality.csv	csv	November 10, 2025, 09:43:28 (UTC+05:30)	1.4 KB	Standard

Actions available for the object include: Copy S3 URI, Copy URL, Download, Open, Delete, Actions (dropdown), Create folder, and Upload.

The bottom of the screen shows the Windows taskbar with icons for CloudShell, Feedback, Search, MobaXterm\_Personal\_25.3 - Shortcut, and various system status indicators like battery level, signal strength, and volume.

The screenshot shows the AWS S3 console interface. The top navigation bar includes tabs for Course Inv, Launch AV, Query edit, s3-daily-re, CloudWatch, logs/2025/, Subscriptions, Redshift, and a plus sign. The main URL in the address bar is [us-east-1.console.aws.amazon.com/s3/object/student-data-lake-34?region=us-east-1&prefix=logs/2025/11/s3\\_athena\\_logs.csv](https://us-east-1.console.aws.amazon.com/s3/object/student-data-lake-34?region=us-east-1&prefix=logs/2025/11/s3_athena_logs.csv). The account ID is 5673-2992-3344, and the user is vclabs/user4548998-rohhamrutkar2525@gmail.com. The region is United States (N. Virginia). The breadcrumb navigation shows the path: Amazon S3 > Buckets > student-data-lake-34 > logs/ > 2025/ > 11/ > s3\_athena\_logs.csv. The object name is s3\_athena\_logs.csv, with an Info button. On the right, there are buttons for Copy S3 URI, Download, Open, and Object actions. Below the object name, there are tabs for Properties (selected), Permissions, and Versions. The Object overview section contains the following details:

Property	Value
Owner	awslabsc0w6299585t1694700763
AWS Region	US East (N. Virginia) us-east-1
Last modified	November 10, 2025, 08:58:41 (UTC+05:30)
Size	2.3 KB
Type	CSV
Key	logs/2025/11/s3_athena_logs.csv
S3 URI	<a href="s3://student-data-lake-34/logs/2025/11/s3_athena_logs.csv">s3://student-data-lake-34/logs/2025/11/s3_athena_logs.csv</a>
Amazon Resource Name (ARN)	<a href="arn:aws:s3:::student-data-lake-34/logs/2025/11/s3_athena_logs.csv">arn:aws:s3:::student-data-lake-34/logs/2025/11/s3_athena_logs.csv</a>
Entity tag (Etag)	<a href="#">6b517304ad536a9296497d8306251921</a>
Object URL	<a href="https://student-data-lake-34.s3.us-east-1.amazonaws.com/logs/2025/11/s3_athena_logs.csv">https://student-data-lake-34.s3.us-east-1.amazonaws.com/logs/2025/11/s3_athena_logs.csv</a>

## Object management overview

The following bucket properties and object management configurations impact the behavior of this object.

## Bucket properties

## Management configurations

Course Inv | Launch AW | Query edit | lambda-da | CloudWat | student-d | Subscriptio | Redshift | +

us-east-1.console.aws.amazon.com/lambda/home?region=us-east-1#/functions/lambda-data-quality-check?tab=configure

aws | Search [Alt+S] | Account ID: 5673-2992-3344  
voldbs/user4548998=rohamrutkar2525@gmail.com

Lambda > Functions > lambda-data-quality-check

## lambda-data-quality-check

Throttle | Copy ARN | Actions ▾

Export to Infrastructure Composer | Download ▾

▼ Function overview [Info](#)

[Diagram](#) | [Template](#)

**lambda-data-quality-check**

Layers (0)

**S3**

+ Add destination

+ Add trigger

Description: An Amazon S3 trigger that retrieves metadata for the object that has been updated.

Last modified: 59 minutes ago

Function ARN: arn:aws:lambda:us-east-1:567329923344:function:lambda-data-quality-check

Function URL | [Info](#)

Code | Test | Monitor | **Configuration** | Aliases | Versions

General configuration [Edit](#)

**General configuration** [Info](#)

Description: An Amazon S3 trigger that retrieves metadata for the object that has been updated.

Memory: 128 MB | Ephemeral storage: 512 MB

Triggers

Permissions

Destinations

CloudShell | Feedback | © 2025, Amazon Web Services, Inc. or its affiliates. | Privacy | Terms | Cookie preferences

Search | Links | MobaXterm\_Personal\_25.3 - Shortcut | Chrome | Edge | YouTube | File Explorer | Task View | Network | Control Panel | Help & Support | Volume | Battery | ENG | 10:41 | Settings

Course Inv | Launch AW | Query edit | lambda-da | CloudWatc | student-d | Subscriptio | Redshift | +

us-east-1.console.aws.amazon.com/lambda/home?region=us-east-1#/functions/lambda-data-quality-check?tab=code

aws | Search [Alt+S] | Account ID: 5673-2992-3344 | vodabs/user4548998=rohamrutkar2525@gmail.com

Lambda > Functions > lambda-data-quality-check

**Code source** Info

**EXPLORER**

LAMBDA-DATA-QUALITY-CHECK

- lambda\_function.py

DEPLOY

- Deploy (Ctrl+Shift+U)
- Test (Ctrl+Shift+I)

TEST EVENTS [NONE SELECTED]

+ Create new test event

**lambda\_function.py**

```
import boto3
import time
import traceback
from datetime import datetime

athena = boto3.client('athena')
sns = boto3.client('sns')

# Configuration
ATHENA_DATABASE = 'sales_db'
ATHENA_OUTPUT = 's3://student-data-lake-34/athena-results/'
SNS_TOPIC_ARN = 'arn:aws:sns:us-east-1:567329923344:data-quality-alerts'

def lambda_handler(event, context):
    start_time = time.time()
    try:
        # Extract file details from S3 event
        record = event['Records'][0]
        bucket = record['s3']['bucket']['name']
        key = record['s3']['object']['key']
        print(f"Triggered by file: {bucket}/{key}")

        # Athena query to find invalid rows
        query = f"""
            SELECT COUNT(*) AS bad_records
            FROM {ATHENA_DATABASE}.orders_raw
        """
        response = athena.start_query_execution(
            QueryString=query,
            QueryExecutionContext={'Database': ATHENA_DATABASE},
            ResultConfiguration={'OutputLocation': ATHENA_OUTPUT}
        )
        execution_id = response['QueryExecutionId']

        # Wait for the query to complete
        while True:
            status = athena.get_query_results(QueryExecutionId=execution_id)
            if status['Status']['State'] == 'SUCCEEDED':
                break
            time.sleep(10)

        # Get the results
        results = athena.get_query_results(QueryExecutionId=execution_id)
        total_bad_records = int(results['ResultSet']['Rows'][0]['C1'][0])
        print(f"Total bad records: {total_bad_records}")

        # Publish an SNS message
        sns.publish(TopicArn=SNS_TOPIC_ARN, Message=f"Bad records found: {total_bad_records} in {bucket}/{key}")
    except Exception as e:
        print(f"Error: {e}")
        sns.publish(TopicArn=SNS_TOPIC_ARN, Message=f"Lambda function failed: {e}")
    finally:
        end_time = time.time()
        duration = end_time - start_time
        print(f"Execution time: {duration:.2f} seconds")
```

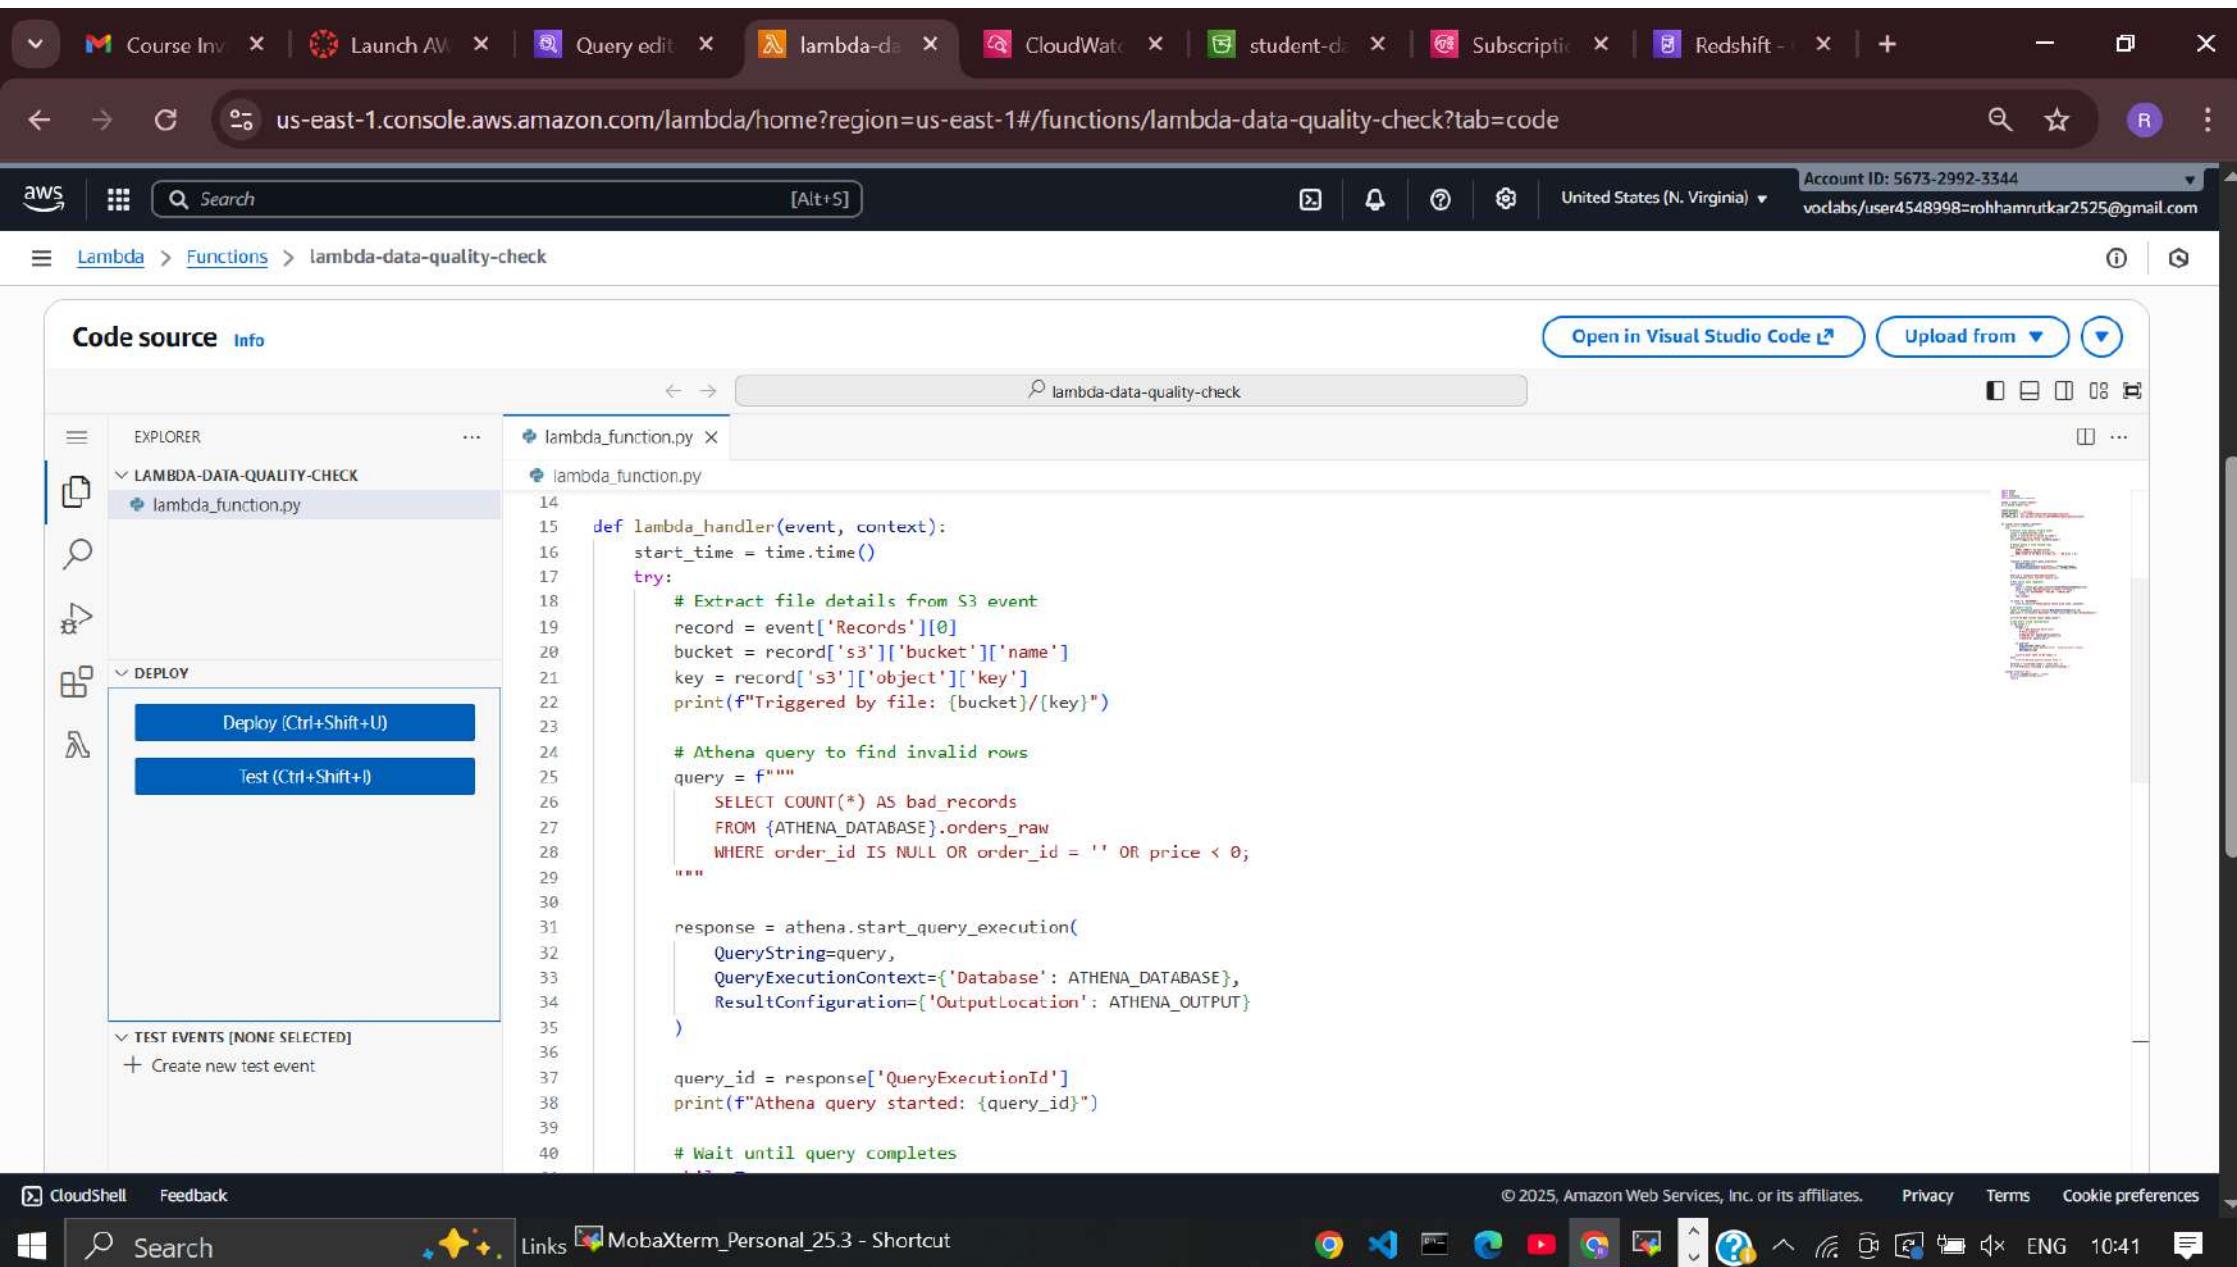
Open in Visual Studio Code | Upload from | ...

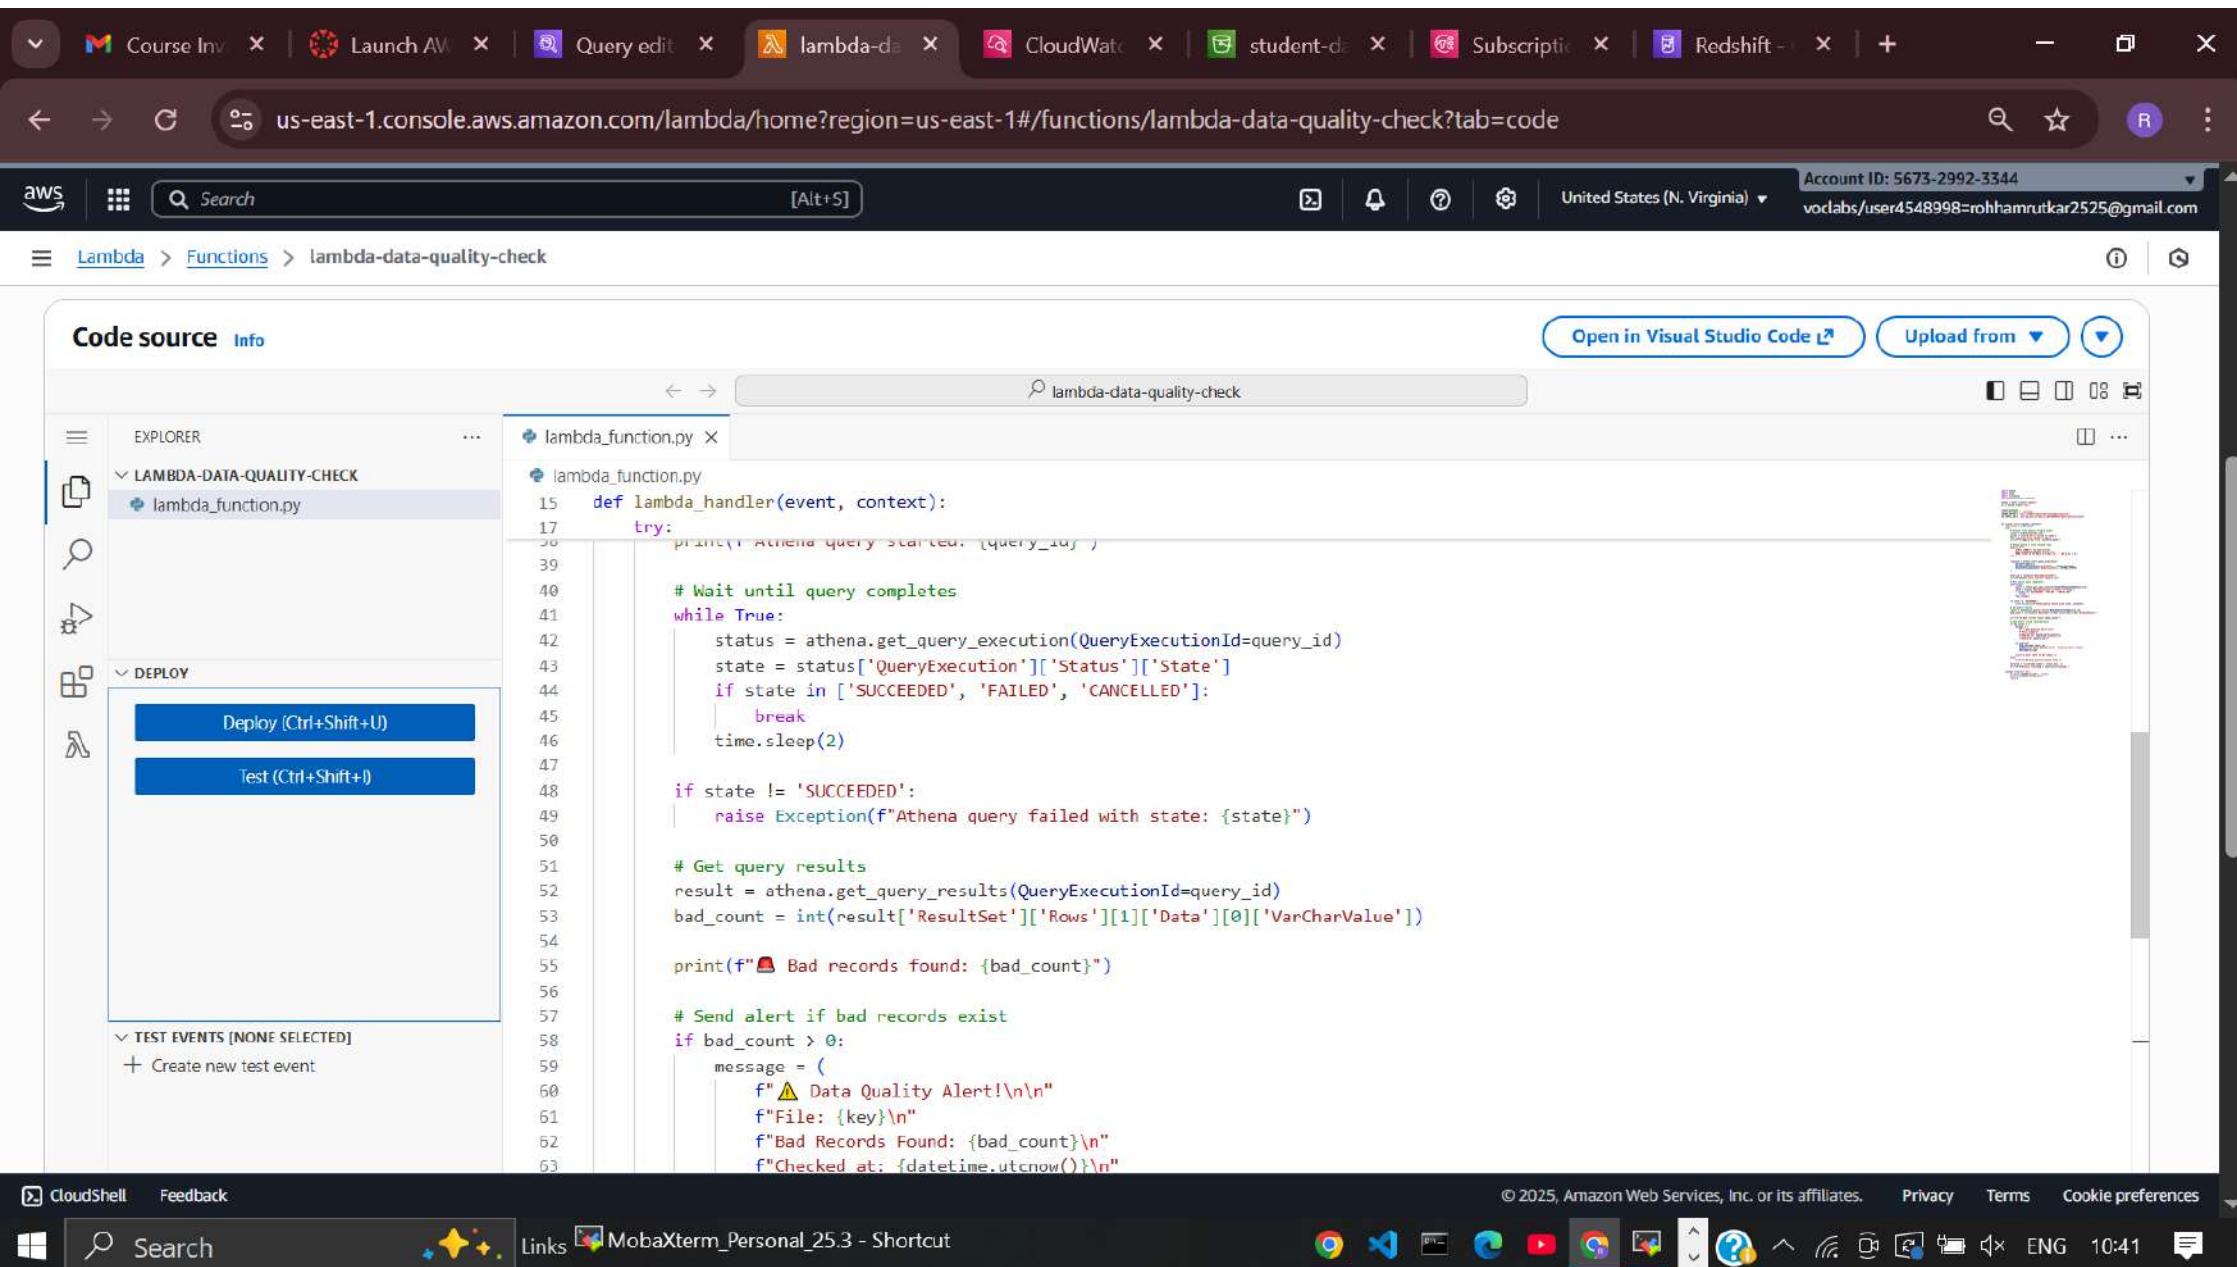
CloudShell Feedback

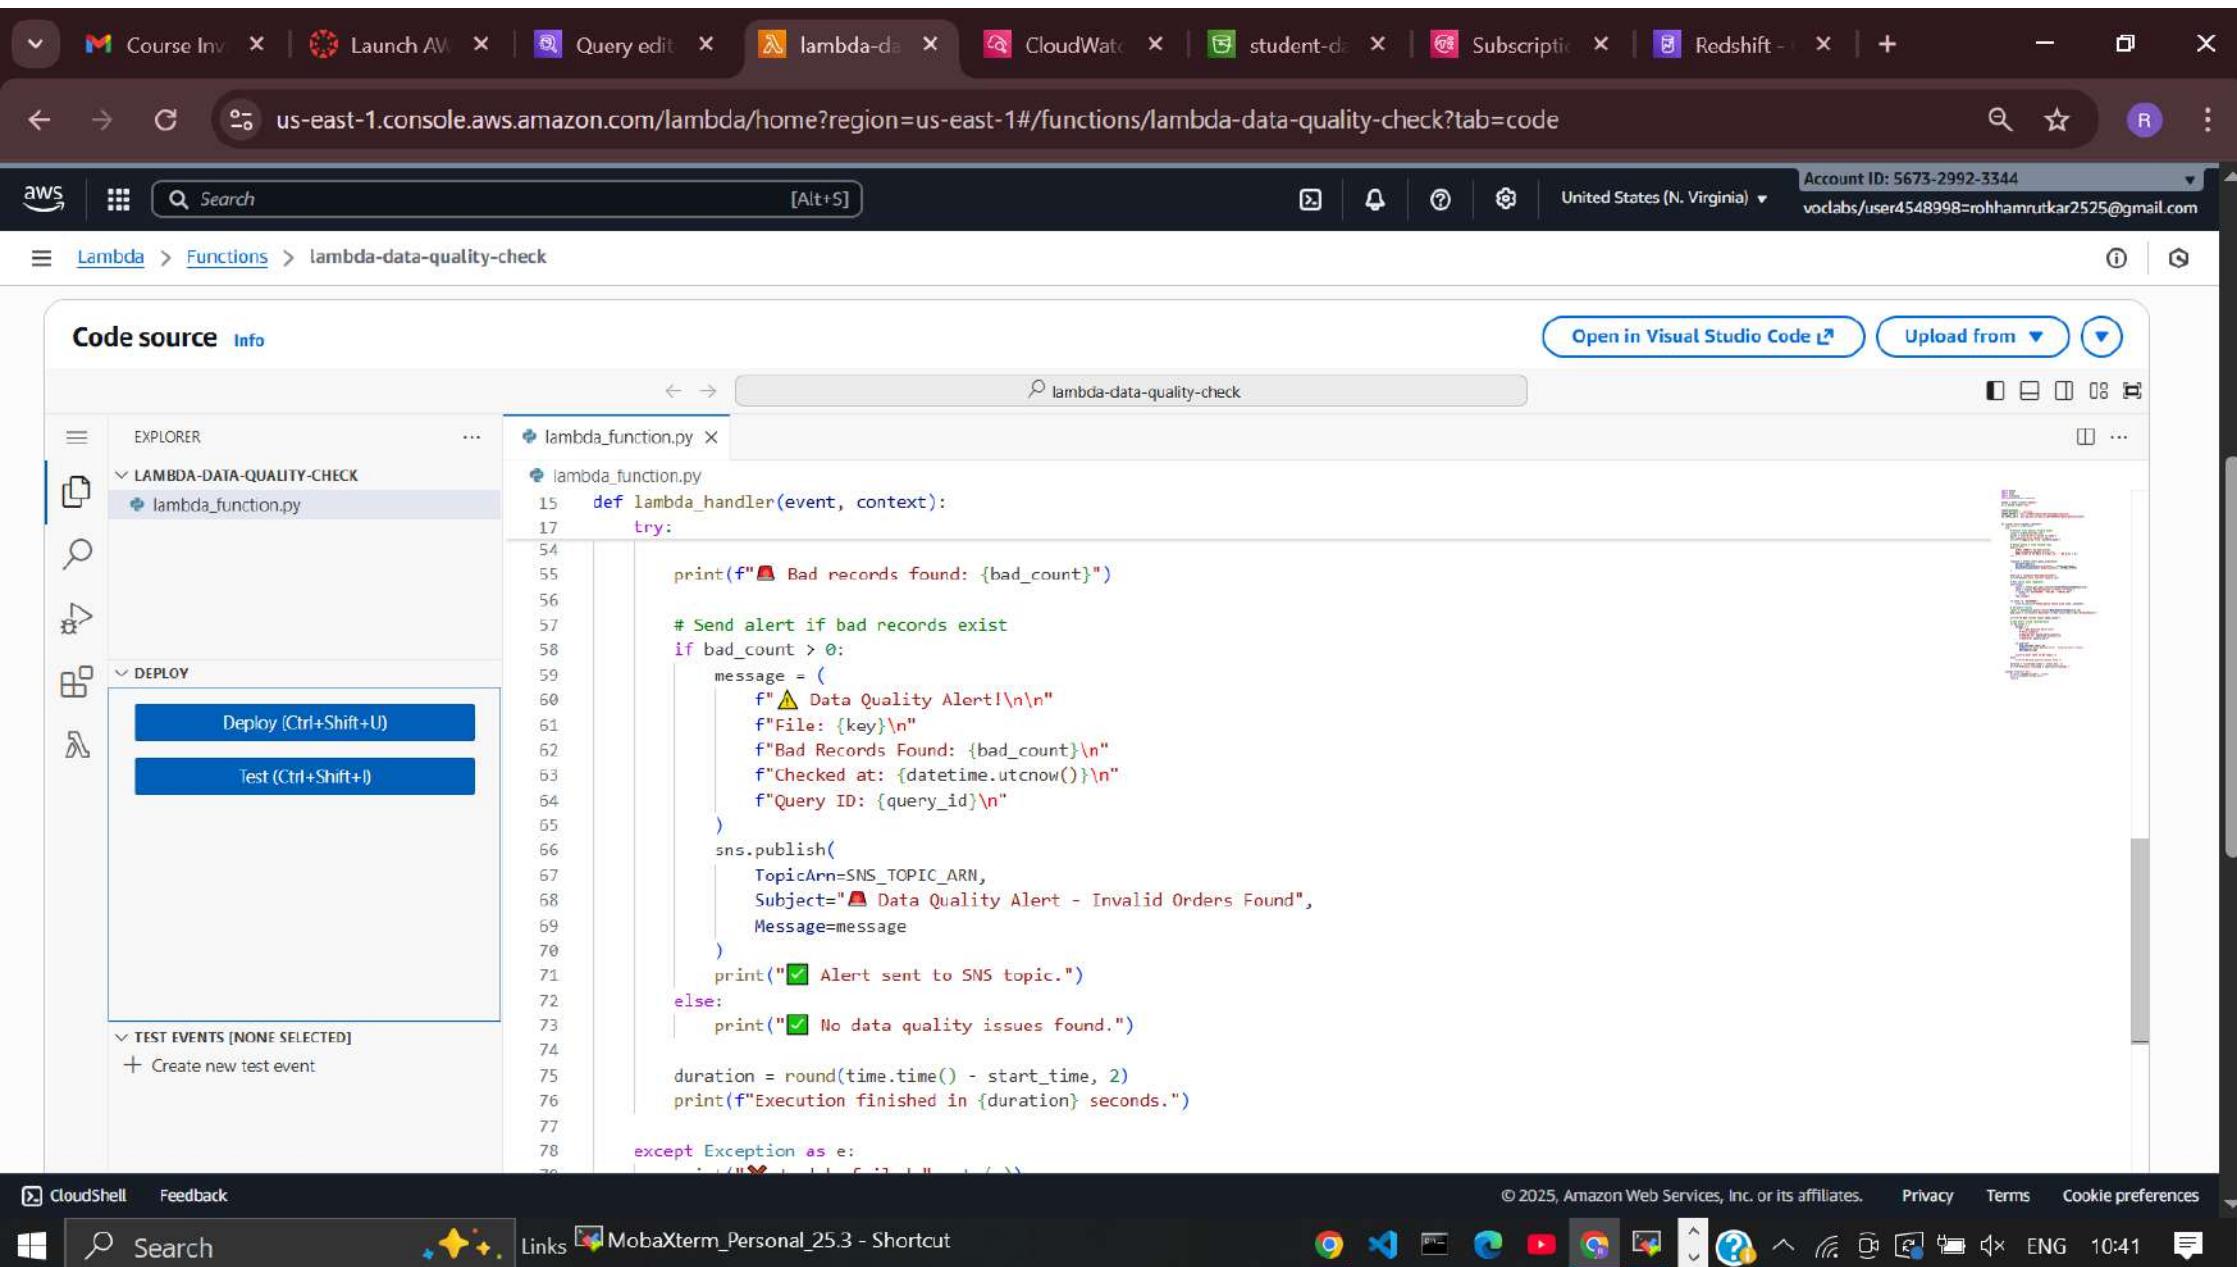
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Search Links MobaXterm\_Personal\_25.3 - Shortcut

Chrome Edge File Explorer YouTube Google Photos Question Mark Network Sound ENG 10:41







Course Inv | Launch AW | Query edit | lambda-da | CloudWatc | student-d | Subscriptio | Redshift | +

us-east-1.console.aws.amazon.com/lambda/home?region=us-east-1#/functions/lambda-data-quality-check?tab=code

aws | Search [Alt+S] | Account ID: 5673-2992-3344  
vodabs/user4548998=rohamrutkar2525@gmail.com

Lambda > Functions > lambda-data-quality-check

**Code source** Info Open in Visual Studio Code Upload from

EXPLORER LAMBDA-DATA-QUALITY-CHECK lambda\_function.py

DEPLOY Deploy (Ctrl+Shift+U) Test (Ctrl+Shift+I)

TEST EVENTS [NONE SELECTED] Create new test event

lambda\_function.py

```
15 def lambda_handler(event, context):
16     try:
17         if bad_count > 0:
18             message = (
19                 f"Checked at: {datetime.utcnow()}\n"
20                 f"Query ID: {query_id}\n"
21             )
22             sns.publish(
23                 TopicArn=SNS_TOPIC_ARN,
24                 Subject="🔴 Data Quality Alert - Invalid Orders Found",
25                 Message=message
26             )
27             print("✅ Alert sent to SNS topic.")
28         else:
29             print("✅ No data quality issues found.")
30
31         duration = round(time.time() - start_time, 2)
32         print(f"Execution finished in {duration} seconds.")
33
34     except Exception as e:
35         print("✖ Lambda failed:", str(e))
36         print(traceback.format_exc())
37         raise e
```

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Search Links MobaXterm\_Personal\_25.3 - Shortcut

Chrome Edge Microsoft Store YouTube Google Photos File Explorer Task View Volume Control Network Connection Sound Volume ENG 10:41

The screenshot shows the AWS CloudWatch interface for a specific log group. The left sidebar navigation includes sections for AI Operations, Alarms, Logs (with Log groups selected), Metrics, Application Signals (APM), GenAI Observability, Network Monitoring, and Insights. The main content area displays the log group details for '/aws/lambda/lambda-data-quality-check'. The ARN is listed as arn:aws:logs:us-east-1:567329923344:log-group:/aws/lambda/lambda-data-quality-check:\*. The log group was created 58 minutes ago and has a retention policy of Never expire. It contains 0 stored bytes. There are no metric filters, subscription filters, or Contributor Insights rules. Data protection settings are listed as '-' for sensitive data count and custom field indexes. A Transformer is also present. The Log streams tab is active, showing one log stream. The stream can be deleted, created, or searched. Filter options include exact match and show expired. The last event time is indicated as 1.

Course | X | Launch | X | Query | X | lambda | CloudW | X | CloudW | X | student | X | Subscri | X | Redshift | X | +

us-east-1.console.aws.amazon.com/cloudwatch/home?region=us-east-1#logsV2.log-groups/log-group/\$252Faws\$252Flambda\$252Flambda-data-quality...

aws Search [Alt+S] Account ID: 5673-2992-3344  
CloudWatch Log groups /aws/lambda/lambda-data-quality-check 2025/11/10/[LATEST]f23d91587ed54cda9736f18475295d65 United States (N. Virginia) vorlabs/user4548998=rohamrutkar2525@gmail.com

CloudWatch Log events

Favorites and recents Dashboards AI Operations New Alarms ▲ 0 ○ 0 ○ 0 Logs Log groups Log Anomalies Live Tail Logs Insights Contributor Insights Metrics New Application Signals New (APM) GenAI Observability Network Monitoring Insights Settings

No older events at this moment. [Retry](#)

INIT\_START Runtime Version: python:3.12.v93 Runtime Version ARN: arn:aws:lambda:us-east-1:runtime:c010741a32d362a2a9a445adf0a7418ce4fac7ec907bcf1371249e79b7ece18a

START RequestId: 6bfe0652-de6e-4987-bf5d-5aa2436a5eea Version: \$LATEST

Triggered by file: student-data-lake-34/incoming/pipeline\_data\_quality.csv

Athena query started: 3dce17b0-238b-4009-8f16-7b77ddbd58ae

Bad records found: 8

Alert sent to SNS topic.

Execution finished in 2.83 seconds.

END RequestId: 6bfe0652-de6e-4987-bf5d-5aa2436a5eea

REPORT RequestId: 6bfe0652-de6e-4987-bf5d-5aa2436a5eea Duration: 2832.20 ms Billed Duration: 3307 ms Memory Size: 128 MB Max Memory Used: 87 MB Init Duration: 474.29 ms

Loading newer events

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Search Links MobaXterm\_Personal\_25.3 - Shortcut Chrome Edge YouTube File Explorer File Manager Task View Volume Control Power User Settings

The screenshot shows the AWS CloudWatch Log Events interface. The left sidebar navigation includes 'CloudWatch' (selected), 'Favorites and recents', 'Dashboards', 'AI Operations', 'Alarms', 'Logs' (selected), 'Log groups', 'Log Anomalies', 'Live Tail', 'Logs Insights', 'Contributor Insights', 'Metrics', 'Application Signals (APM)', 'GenAI Observability' (New), 'Network Monitoring' (New), and 'Insights'. The main content area is titled 'Log events' and displays a single log entry for a Lambda function execution. The log message is as follows:

```
INIT_START Runtime Version: python:3.12.v93      Runtime Version ARN: arn:aws:lambda:us-east-1::runtime:c010741a32d362a2a0a445adf0a7418ce4fac7ac907bcf1371249e79b7ece18a
START RequestId: 6bfe0652-de6e-4987-bf5d-5aa2436a5eea Version: $LATEST
Triggered by file: student-data-lake-34/incoming/pipeline_data_quality.csv
Athena query started: 3dce17b9-230b-4009-8f16-7b77ddbd58ae
Bad records found: 8
Alert sent to SNS topic.
Execution finished in 2.83 seconds.
END RequestId: 6bfe0652-de6e-4987-bf5d-5aa2436a5eea
REPORT RequestId: 6bfe0652-de6e-4987-bf5d-5aa2436a5eea Duration: 2832.20 ms      Billed Duration: 3307 ms      Memory Size: 128 MB      Max Memory Used: 87 MB      Init Duration: 474.29 ms
```

Below the log message, it says 'No newer events at this moment. Auto retry paused. [Resume](#)'.

